

# Lab 4

*Shan He, Joanna Huang, Tiffany Jaya*

*17 December 2017*

## Introduction

This year, on October 18, 2017, Law Enforcement Leaders urged Attorney General Jeff Sessions to reconsider his stance on reverting back to “overly punitive” approaches of the 1980s and 1990s to reduce crime. Since President Trump believes that America is in the midst of a national crime wave, Sessions thought a more conservative approach of deterrence through arrests, incapacitation through imprisonment, harsh sentencing and higher police per capita would lead to lower crime rates overall. However, police chiefs who have first hand decades of experience on the front lines learned that these tactics are ineffective to reduce crime.

In this paper, we will explore whether the conservative approach to crime effectively reduce crime rates. We began by exploring North Carolina’s crime dataset of 1988 when “overly punitive” approaches of the 1980s and 1990s would have taken place and analyzed the determinants of crime based on the research question: Does the conservative approach of deterrence through arrests, incapacitation through imprisonment, harsh sentencing and higher police per capita lead to lower crime rates? We will list out the limitations of our analysis, including any estimates that suffer from endogeneity bias, and generate policy suggestions based on our findings.

## Exploratory Data Analysis

```
# load the data
data <- read.csv("crime_v2_updated.csv")
# verify that it only contains data from 1988
unique(data$year)

## [1] 88

# list number of counties
length(unique(data$county))

## [1] 90

# list number of western, central, and urban counties
c(sum(data$west == 1), sum(data$central == 1), sum(data$urban == 1))

## [1] 34 21 8

# list number of western & urban counties and central & urban counties
c(sum(data$west == 1 & data$urban == 1), sum(data$central == 1 & data$urban == 1))

## [1] 5 1

# verify number of missing values
colSums(sapply(data, is.na))

##      X   county   year   crime  probarr  probsen  probconv  avgsen
##      0      0      0      0      0      0      0      0
## police density   tax    west  central   urban  pctmin  wagecon
##      0      0      0      0      0      0      0      0
```

```
##   wagetuc   wagetrd   wagefir   wageser   wagemfg   wagefed   wagesta   wageloc
##         0         0         0         0         0         0         0         0
##      mix     ymale
##         0         0
```

The dataset contains 90 counties from North Carolina, all of which is collected in 1988. Out of the 90 counties, 34 are from western NC (out of which 5 is also urban), 21 are from central NC (out of which 1 is also urban), and 8 are considered urban counties. There are no missing values which will make our analysis easier.

```
summary(data)
```

```
##           X           county           year           crime
## Min.      : 1.00   Min.      : 1.0   Min.      :88   Min.      :0.005533
## 1st Qu.:23.25   1st Qu.: 51.5   1st Qu.:88   1st Qu.:0.020604
## Median :45.50   Median :103.0   Median :88   Median :0.030002
## Mean     :45.50   Mean     :100.6   Mean     :88   Mean     :0.033510
## 3rd Qu.:67.75   3rd Qu.:150.5   3rd Qu.:88   3rd Qu.:0.040249
## Max.     :90.00   Max.     :197.0   Max.     :88   Max.     :0.098966
##   probarr      probsen      probconv      avgsen
## Min.      :0.1500   Min.      :0.09277   Min.      :0.06838   Min.      : 5.380
## 1st Qu.:0.3642   1st Qu.:0.20495   1st Qu.:0.34422   1st Qu.: 7.375
## Median :0.4222   Median :0.27146   Median :0.45170   Median : 9.110
## Mean     :0.4106   Mean     :0.29524   Mean     :0.55086   Mean     : 9.689
## 3rd Qu.:0.4576   3rd Qu.:0.34487   3rd Qu.:0.58513   3rd Qu.:11.465
## Max.     :0.6000   Max.     :1.09091   Max.     :2.12121   Max.     :20.700
##   police      density      tax      west
## Min.      :0.0007459   Min.      :0.2034   Min.      : 25.69   Min.      :0.0000
## 1st Qu.:0.0012378   1st Qu.:0.5472   1st Qu.: 30.73   1st Qu.:0.0000
## Median :0.0014897   Median :0.9792   Median : 34.92   Median :0.0000
## Mean     :0.0017080   Mean     :1.4379   Mean     : 38.16   Mean     :0.3778
## 3rd Qu.:0.0018856   3rd Qu.:1.5693   3rd Qu.: 41.01   3rd Qu.:1.0000
## Max.     :0.0090543   Max.     :8.8277   Max.     :119.76   Max.     :1.0000
##   central      urban      pctmin      wagecon
## Min.      :0.0000   Min.      :0.00000   Min.      : 1.284   Min.      :193.6
## 1st Qu.:0.0000   1st Qu.:0.00000   1st Qu.:10.024   1st Qu.:250.8
## Median :0.0000   Median :0.00000   Median :24.852   Median :281.2
## Mean     :0.2333   Mean     :0.08889   Mean     :25.713   Mean     :285.4
## 3rd Qu.:0.0000   3rd Qu.:0.00000   3rd Qu.:38.183   3rd Qu.:315.0
## Max.     :1.0000   Max.     :1.00000   Max.     :64.348   Max.     :436.8
##   wagetuc      wagetrd      wagefir      wageser
## Min.      :187.6   Min.      :154.2   Min.      :170.9   Min.      : 133.0
## 1st Qu.:374.3   1st Qu.:190.7   1st Qu.:285.6   1st Qu.: 229.3
## Median :404.8   Median :203.0   Median :317.1   Median : 253.1
## Mean     :410.9   Mean     :210.9   Mean     :321.6   Mean     : 275.3
## 3rd Qu.:440.7   3rd Qu.:224.3   3rd Qu.:342.6   3rd Qu.: 277.6
## Max.     :613.2   Max.     :354.7   Max.     :509.5   Max.     :2177.1
##   wagemfg      wagefed      wagesta      wageloc
## Min.      :157.4   Min.      :326.1   Min.      :258.3   Min.      :239.2
## 1st Qu.:288.6   1st Qu.:398.8   1st Qu.:329.3   1st Qu.:297.2
## Median :321.1   Median :448.9   Median :358.4   Median :307.6
## Mean     :336.0   Mean     :442.6   Mean     :357.7   Mean     :312.3
## 3rd Qu.:359.9   3rd Qu.:478.3   3rd Qu.:383.2   3rd Qu.:328.8
## Max.     :646.9   Max.     :598.0   Max.     :499.6   Max.     :388.1
##      mix      ymale
## Min.      :0.01961   Min.      :0.06216
```

```
## 1st Qu.:0.08060 1st Qu.:0.07437
## Median :0.10095 Median :0.07770
## Mean :0.12905 Mean :0.08403
## 3rd Qu.:0.15206 3rd Qu.:0.08352
## Max. :0.46512 Max. :0.24871
```

Most of the variables appear to be within a reasonable range, except for *probarr* and *probconv*, which have probability values greater than 1.

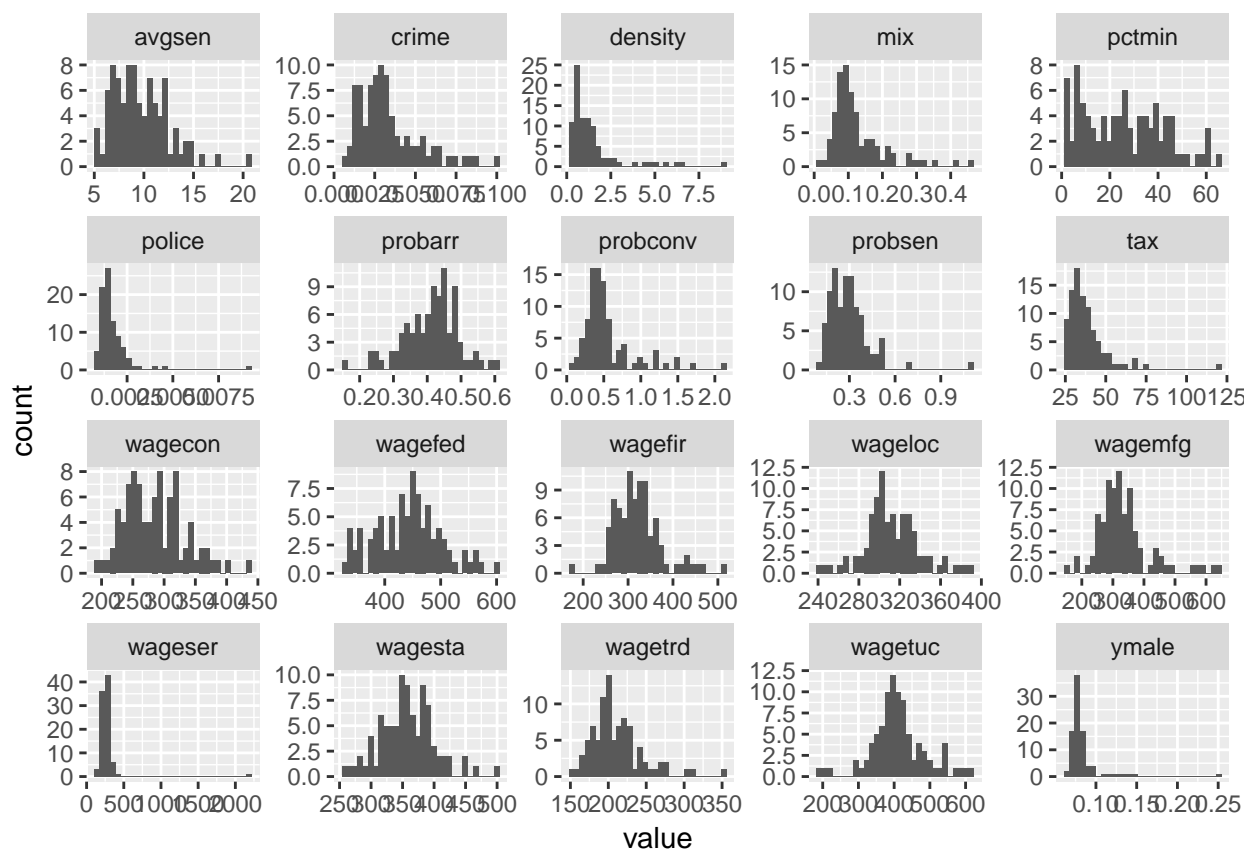
```
# list number of probabilities (probarr, probconv, probsen, mix) that are not in range [0, 1]
c(sum(data$probarr < 0 | 1 < data$probarr), sum(data$probconv < 0 | 1 < data$probconv),
  sum(data$probsen < 0 | 1 < data$probsen), sum(data$mix < 0 | 1 < data$mix))
```

```
## [1] 0 10 1 0
```

*probconv* and *probsen* contain 10 and 1 datapoints respectively that do not conform to the probability assumption. We will take these outliers into consideration when choosing variables for our models.

We then plot each numeric variable in a histogram to see its sample distribution.

```
# plot every variable except X, county, year, west, central, urban
num.data <- data[!(names(data) %in% c("X", "county", "year", "west", "central", "urban"))]
ggplot(gather(num.data), aes(value)) +
  facet_wrap(~key, scales="free") +
  geom_histogram()
```



```
skewness(num.data)
```

```
##      crime      probarr      probsen      probconv      avgсен      police
## 1.28174888 -0.45254022  2.52529596  2.03950599  1.00116340  4.98348795
```

```
##      density      tax      pctmin      wagecon      wagetuc      wagetrd
## 2.65301071 3.29057447 0.36566169 0.60680223 0.06819768 1.46120657
##      wagefir      wageser      wagemfg      wagefed      wagesta      wageloc
## 0.82063145 8.69918165 1.42253166 0.13223761 0.36236826 0.29513808
##      mix      ymale
## 1.91657046 4.56069074
```

Most of the sample distributions appear to be positively skewed. When choosing the variables for our regression models, we will consider logarithmic transformations if the interpretations make sense.

From the histograms, we also see several notable outliers. We are under the impression that a county which has an outlier in one variable will likely have an outlier in another variable. For this reason, we have listed counties which have repeated outliers when we iterate through the entire numeric variables.

```
# iterate through each numeric variable and list the outlier counties and their respective frequency
county.ids <- c()
for(var in num.data) {
  var.out <- boxplot.stats(var)$out
  county.ids <- c(county.ids, data[var %in% var.out, ]$county)
}
table(county.ids)
```

```
## county.ids
##  1  3  5  7 11 19 35 39 49 51 53 55 63 67 69 71 79 81
##  1  1  1  1  2  4  2  2  1  3  1  3  5  1  3  2  1  2
## 85 87 93 99 105 111 113 115 119 123 127 129 131 133 135 137 139 143
##  1  1  1  2  1  1  1  5 10  1  2  3  1  1  2  2  1  2
## 147 149 169 173 175 181 183 185 187 189 195 197
##  1  1  1  4  1  2  4  2  1  1  2  1
```

```
# list the most extreme outlier
outlier(num.data)
```

```
##      crime      probarr      probsen      probconv      avgsgen
## 0.09896590 0.15000001 1.09090996 2.12121010 20.70000076
##      police      density      tax      pctmin      wagecon
## 0.00905433 8.82765198 119.76145170 64.34819794 436.76663210
##      wagetuc      wagetrd      wagefir      wageser      wagemfg
## 187.61726380 354.67611690 509.46551510 2177.06811500 646.84997560
##      wagefed      wagesta      wageloc      mix      ymale
## 597.95001220 499.58999630 388.08999630 0.46511629 0.24871162
```

One outlier that is interesting to note is the weekly wage in the service industry for county with id 185, \$2177.10.

```
summary(data$wageser)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 133.0   229.3   253.1   275.3   277.6  2177.1
```

It is approximately eight times higher than the median. We do not know if the value is inputted incorrectly or if the county in general is making a weekly wage of \$2177.10 in the service industry.

## Research Question

James Q. Wilson and George Kelling's "broken windows theory" in 1982 led to a nation-wide movement for stricter crime-fighting policies between the 1980s and 1990s. The theory states:

*if the first broken window in a building is not repaired, then people who like breaking windows will assume that no one cares about the building and more windows will be broken. Soon the building will have no windows...*





The belief was that by adopting a zero tolerance approach that enforced even the lowest level offenses, crime rates would subsequently go down. While New York City notably enforced this more stringent approach, San Francisco went the opposite direction of less strident law enforcement policies that reduced arrests, prosecutions and incarceration rates. Both sides experienced considerable declines in crime rates. Thus we hope to test the “broken windows theory” for the counties of South Carolina in 1987 and answer the question: Does the conservative approach of deterrence through arrests, incapacitation through imprisonment, harsh sentencing and higher police per capita lead to lower crime rates?

## Model 1: only the explanatory variables of key interest

Based on the research question, our initial proposed model will include *crime* as the dependent variable and all variables related to stricter law enforcement policies: *probarr*, *probconv*, *probsen*, *avgpen*, and *police* as independent variables. Assuming the “broken windows theory” is valid, we expect generally negative coefficients for all variables.

Given that the histogram of *crime* has a significant positive skew, we noted a log transformation may be suitable since its values are non-zero and positive. The same can be said about the independent variable *police* where its histogram is positively skewed and its values are non-zero and positive.

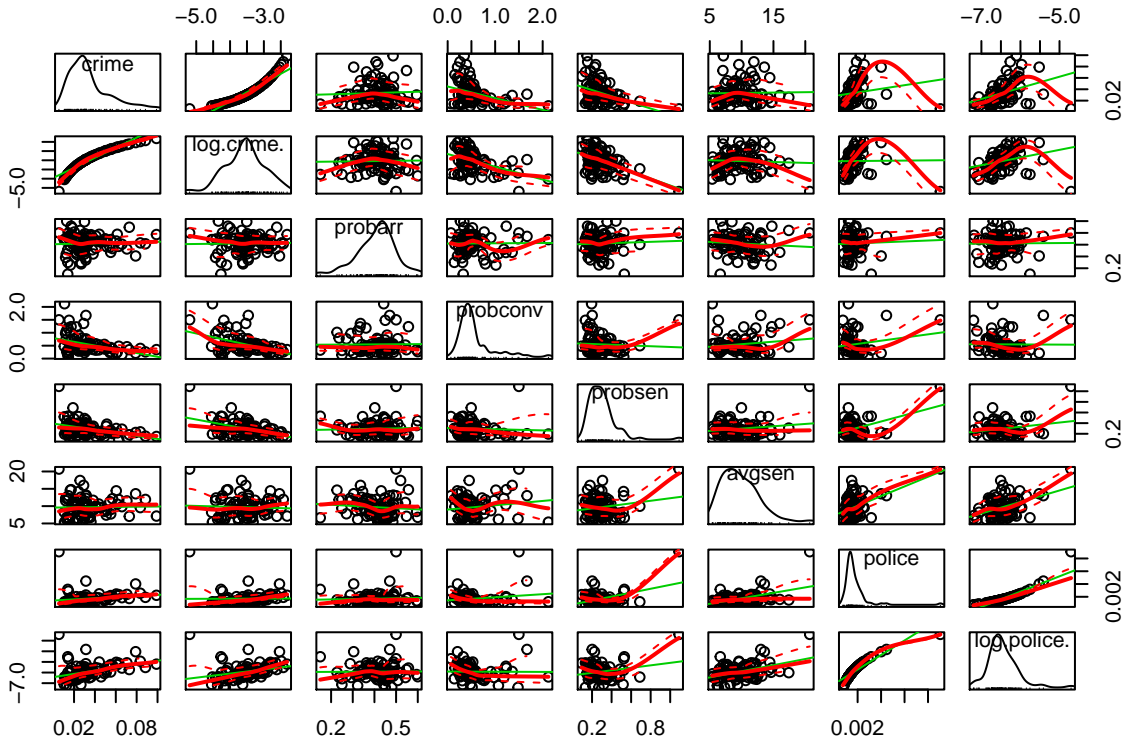
```
# before and after log transformation
hist(data$crime); hist(log(data$crime))
hist(data$police); hist(log(data$police))
```



Though *probarr*, *probconv*, and *probsen* are positively skewed as well, we decided against taking the log of these variables because log transformations can make values between 0 and 1 more extreme. We also kept *avgsen* as is for easier interpretation.

Next, we want to check the relationships between the chosen independent variables and our dependent variable, before and after transformations. We want to ensure that we did not deviate any straight-line relationships between the independent variables and the dependent variable using the transformation.

```
scatterplotMatrix(~ crime + log(crime) +
                  probarr + probconv + probsen + avgsen + police + log(police), data = data)
```



As we can see from the scatterplot matrix, it does not appear that the transformation drastically changed the relationship.

Lastly, based on the exploratory data analysis, we should be careful when considering *probconv* and *probsen* as variables in the model with 10 and 1 datapoints respectively that have probabilities greater than 1. *probconv* is proxied by the ratio of convictions to arrest while *probsen* is proxied by the proportion of total convictions resulting in prison sentences. Although it is unlikely that an individual can be convicted without an arrest or sentenced without a conviction, we cannot rule out the possibility. Both of these variables are important in answering our research question and removing them will result in an omitted variable bias as we will demonstrate below.

Assuming we started out with a base model without *probconv* and *probsen*, we wanted to see what effects *probconv* and *probsen* respectively have on the other explanatory variables when we add them individually to the base model. We looked at the printout of their respective model coefficients to understand the effects. Based on the research question, we expect that higher conviction and higher sentencing will result in lower crime rate. And since the relationship of *probconv* and *probsen* are positive with the other explanatory variables as demonstrated by the correlation matrix, we expect negative bias overall.

```

# demonstrate that probconv and probsen individually have positive relationship
# with the other explanatory variables: probarr, avgsten, police
ind.vars <- subset(data, select= c("probarr", "probconv", "probsen", "avgsten", "police"))
cor(ind.vars, ind.vars)

##          probarr    probconv    probsen    avgsten    police
## probarr    1.00000000  0.01102265  0.04583324 -0.09468083  0.04820783
## probconv    0.01102265  1.00000000 -0.05579621  0.15585232  0.17186514
## probsen     0.04583324 -0.05579621  1.00000000  0.17869425  0.42596480
## avgsten    -0.09468083  0.15585232  0.17869425  1.00000000  0.48815230
## police      0.04820783  0.17186514  0.42596480  0.48815230  1.00000000

# test omitted variable bias by first creating a base model and a model for each omitted variable
m1.base <- lm(crime ~ probarr + avgsten + police, data=data)
m1.probconv <- lm(crime ~ probarr + probconv + avgsten + police, data=data)
m1.probsen <- lm(crime ~ probarr + probsen + avgsten + police, data=data)
# print out the model coefficients
(coef.base <- coeftest(m1.base, vcov=vcovHC))

##
## t test of coefficients:
##
##          Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.02885462  0.02976205  0.9695  0.3350
## probarr      0.00725067  0.03104952  0.2335  0.8159
## avgsten     -0.00050917  0.00082262 -0.6190  0.5376
## police      3.87085953 11.68258001  0.3313  0.7412

(coef.probconv <- coeftest(m1.probconv, vcov=vcovHC))

##
## t test of coefficients:
##
##          Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.03638722  0.02803944  1.2977 0.197896
## probarr      0.00866278  0.03144264  0.2755 0.783593
## probconv    -0.02265313  0.00764791 -2.9620 0.003962 **
## avgsten     -0.00023551  0.00082518 -0.2854 0.776031
## police      4.87482880  9.63460820  0.5060 0.614187
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(coef.probsen <- coeftest(m1.probsen, vcov=vcovHC))

##
## t test of coefficients:
##
##          Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.04399311  0.02316845  1.8988 0.060978 .
## probarr      0.01010303  0.02266319  0.4458 0.656882
## probsen     -0.07884692  0.02817947 -2.7980 0.006359 **
## avgsten     -0.00064360  0.00074499 -0.8639 0.390070
## police      8.71360915  6.17379949  1.4114 0.161781
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Looking at the coefficients, it does appear that higher conviction (*probconv*) and higher sentencing (*probsen*)

result in lower crime rate (*crime*) as seen by their negative sign in their respective coefficient. We also note that *probconv* and *probsen* are statistically significant when added to the base model. It appears that there is a negative omitted variable bias. For this reason, it would be best to include *probconv* and *probsen* in our Model 1 proposal.

As we will later discuss in section “Discussion of Causality”, if the outliers happen to be a measurement error, it will result in our model being confounded by bias. Although that might be the case, there is also a likelihood that the measurement is valid, and we have demonstrated that not including *probconv* and *probsen* will most likely confound our model with omitted variable bias.

Hence, we propose our first model as follows which contains all explanatory variables of key interest:

$$\log(\text{crime}) = \beta_0 + \beta_1 \cdot \text{probarr} + \beta_2 \cdot \text{probconv} + \beta_3 \cdot \text{probsen} + \beta_4 \cdot \text{avgsen} + \beta_5 \cdot \log(\text{police}) + u$$

We will now run the model and test the validity of the 6 CLM assumptions to ensure that the OLS estimators are consistent, normally distributed, and BLUE (best linear unbiased estimator).

```
m1 <- lm(log(crime) ~ probarr + probconv + probsen + avgsen + log(police), data=data)
```

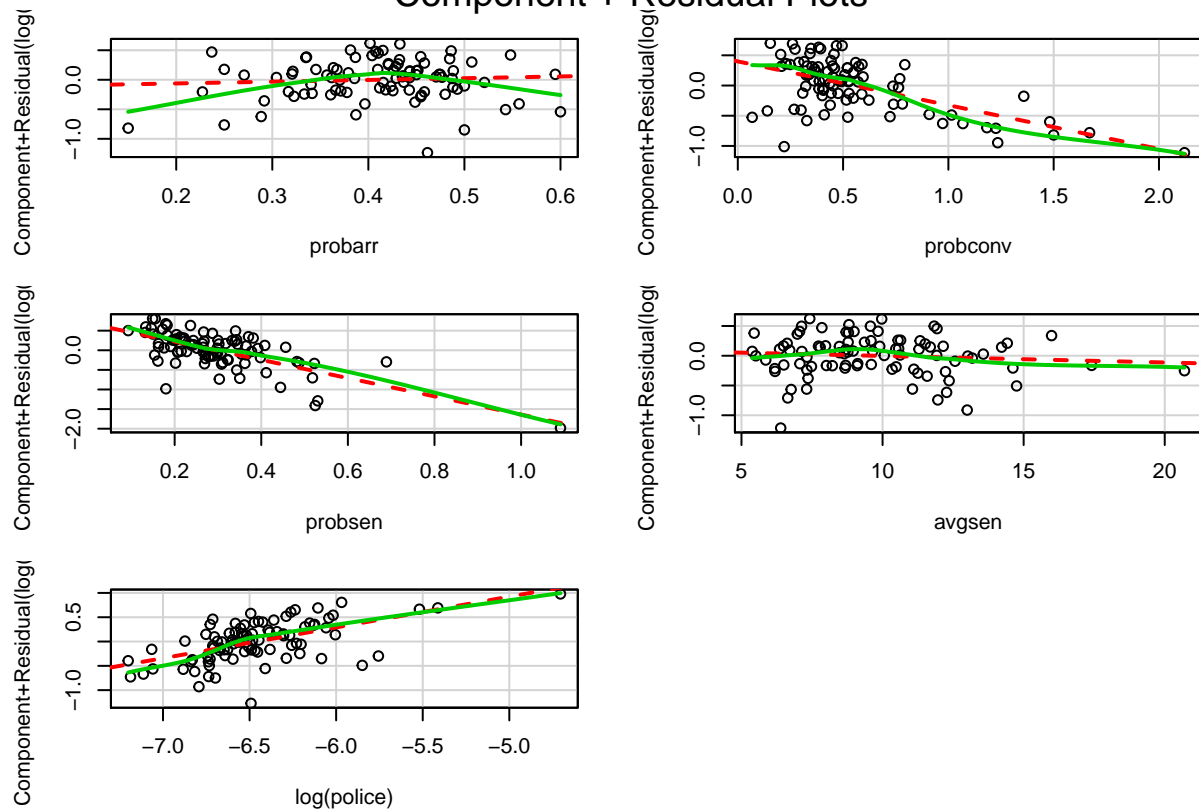
## CLM 1 - A linear model

The model is specified such that the dependent variable is a linear function of the explanatory variables. As shown in the scatterplot matrix above, all of the dependent variables in the model seem to have a linear relationship with the independent variable *log(crime)*. We can verify further the linearity of the relationship using either component+residual plots (also called partial-residual plots) or the CERES plots. We have decided to do the former and note that for the most part, the relationships appear linear.

```
# verify linearity of relationships using component+residual plots
crPlots(m1)
```



## Component + Residual Plots



## CLM 2 - Random Sampling

We do not know how the survey is collected. We assume that the variables are representative of the entire population distribution since the counties are subsets of North Carolina. There is nothing we can do to correct this, so we note this as a potential weakness in the analysis.

## CLM 3 - Multicollinearity

As a quick test of the multicollinearity condition, we check the correlation of the explanatory variables and their Variance Inflation Factors (VIF):

```
# correlation matrix of explanatory variables
data$log.police <- log(data$police)
cor(data.matrix(
  subset(data, select=c("probarr", "probconv", "probsen", "avgsten", "police", "log.police"))))
```

```
##          probarr    probconv    probsen    avgsten    police
## probarr    1.0000000    0.011022645    0.04583324   -0.09468083    0.04820783
## probconv    0.01102265    1.000000000   -0.05579621    0.15585232    0.17186514
## probsen     0.04583324   -0.055796206    1.00000000    0.17869425    0.42596480
## avgsten    -0.09468083    0.155852319    0.17869425    1.00000000    0.48815230
## police      0.04820783    0.171865142    0.42596480    0.48815230    1.00000000
## log.police  0.01041349   -0.007574593    0.21624362    0.43729326    0.90577332
##          log.police
## probarr    0.010413494
```

```
## probconv -0.007574593
## probsen 0.216243619
## avgsen 0.437293263
## police 0.905773321
## log.police 1.000000000
```

```
# verify VIFs are less than 10
vif(m1)
```

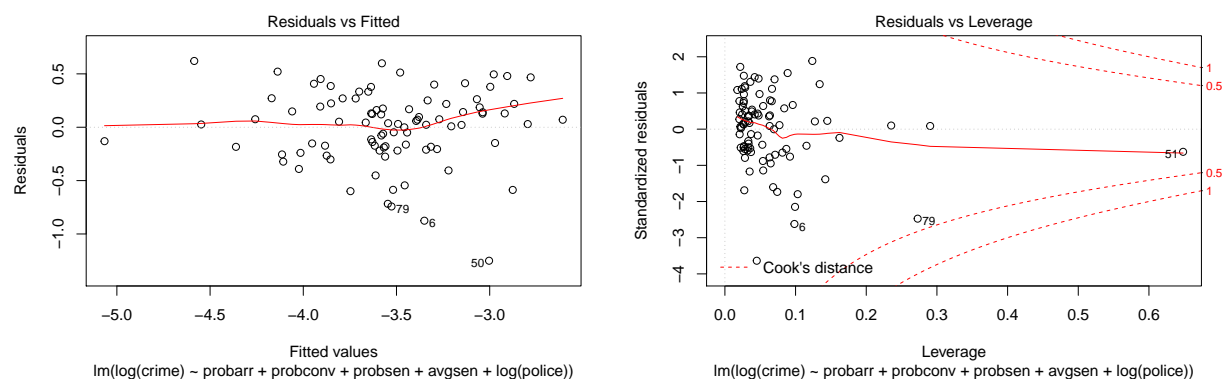
```
##      probarr      probconv      probsen      avgsen log(police)
## 1.016889 1.039388 1.068228 1.310152 1.277425
```

The explanatory variables (*probarr*, *probconv*, *prbpis*, *avgsen*, *log.police*) are not perfectly correlated and the VIFs are low (i.e. less than 10), so there is no perfect multicollinearity of the independent variables.

## CLM 4 – Zero-Conditional Mean

To see whether there is a zero-conditional mean across all  $x$ 's, we will plot the residuals against the fitted values.

```
# plot residual vs fitted plot & residual vs leverage plot
plot(m1, which=c(1, 5))
```



The residual vs fitted plot indicates little evidence that the zero-conditional mean assumption does not hold since the red spline line remains close to zero despite its slight dip and rise at both ends due to fewer observations.

Furthermore, it does not appear that the outliers have undue influence on the model fit. Based on the residual vs leverage plot, none of the outliers have a leverage that exceeds a Cook's distance of 1 on the regression model.

We have also taken a look at the covariances of the independent variables with the residuals to see if the variables we chose are likely to be exogenous.

```
# calculate the covariance for each independent variables with the model's residuals
lapply(subset(data, select=c("probarr", "probconv", "probsen", "avgsen", "log.police")),
       function(var) cov(var, m1$residuals))
```

```
## $probarr
## [1] 0.000000000000000008087166
##
## $probconv
## [1] 0.000000000000000003500816
```



```
## Chisquare = 1.496155    Df = 1    p = 0.2212639
```

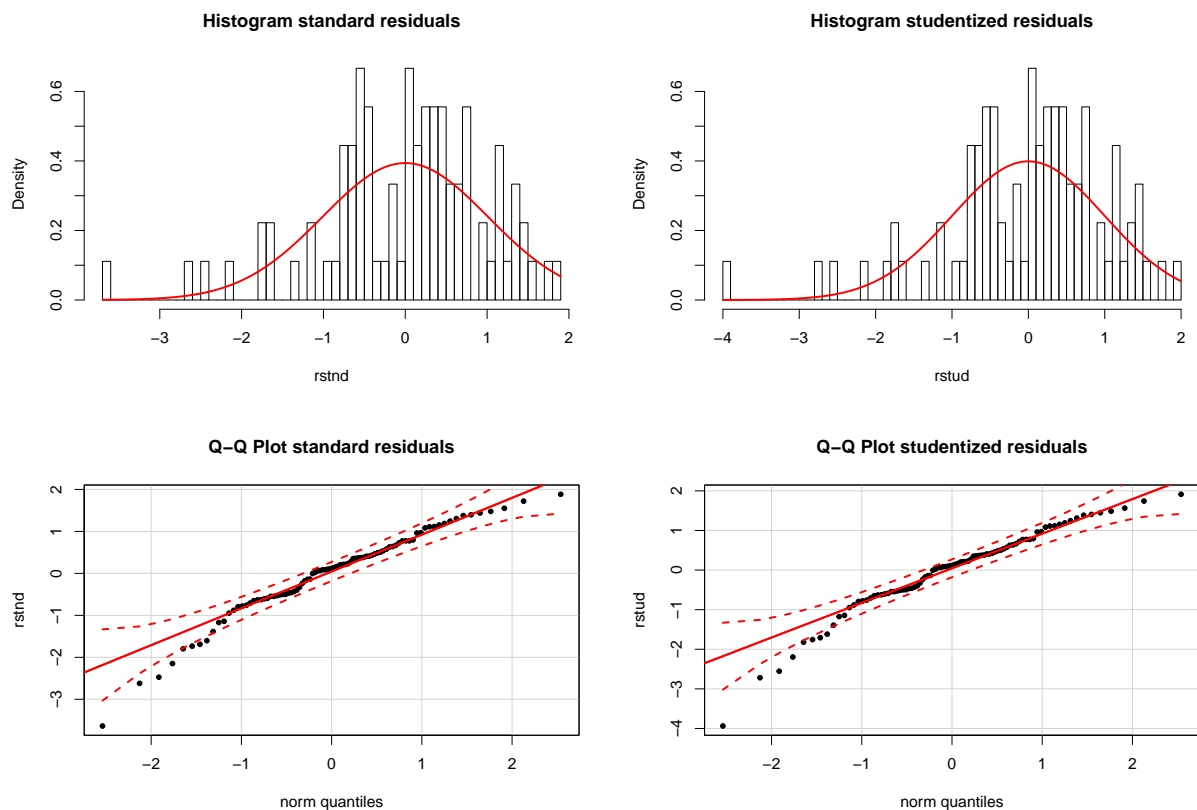
The Score-test also validates this assumption. Since the p-value is statistically not significant, we cannot reject the null hypothesis of constant error variance.

For this reason, the assumption of homoskedasticity is met.

## CLM 6 – Normality of residuals

To determine whether there is normality of the residuals, we looked at the histogram and the Q-Q plot of the residuals and visually observe whether there is normality.

```
# normality of standard residuals
rstnd = rstandard(m1)
hist(rstnd, main="Histogram standard residuals", breaks=50, freq=FALSE)
curve(dnorm(x, mean=0, sd=sd(rstnd)), col="red", lwd=2, add=TRUE)
# normality of studentized residuals
rstud = rstudent(m1)
hist(rstud, main="Histogram studentized residuals", breaks=50, freq=FALSE)
curve(dnorm(x, mean=0, sd=1), col="red", lwd=2, add=TRUE)
# Q-Q plot standard residuals
qqPlot(rstnd, distribution="norm", pch=20, main="Q-Q Plot standard residuals")
qqline(rstnd, col="red", lwd=2)
# Q-Q plot studentized residuals
qqPlot(rstud, distribution="norm", pch=20, main="Q-Q Plot studentized residuals")
qqline(rstud, col="red", lwd=2)
```



The histograms appear to be negatively skewed. The Q-Q plots further supports it with a fat negative tail.

```
#check sample size for model 1
nobs(m1)
```

```
## [1] 90
```

Although the assumption is not met, given the substantial sample size, we can be confident that due to OLS asymptotics the distribution of the residuals will be approximately normal.

Since all six assumptions of the Classical Linear Model are met, we can assume that the OLS estimators are consistent, normally distributed and BLUE.

## Model 2: add covariates that increase accuracy without bias

For Model 2, we decided to include variables that have an indirect impact to crime rate: *density*, *tax*, and *mix*.

We chose *density* based on the theory that the more densely populated an area is, the harder it is for individuals to commit crime, which in turn decreases the crime rate. We assumed that if we included *density* on top of Model 1, its coefficient will be positive since it has a negative relationship with *crime* and a negative relationship with most of the explanatory variables except *police* as shown below. Our thought process goes as follows: because we assume that densely populated area will lower crime rate, it will lower the probability of arrest, conviction and prison sentence and therefore have a negative relationship with them. On the other hand, we assume that an increased number of people per capita will reflect an increased number of police per capita and therefore *density* will have a positive relationship with *police*.


```
# list the correlation between density and model 1's explanatory variables
cor(subset(data, select="density"),
    subset(data, select=c("probarr", "probconv", "probsen", "avgsen", "log.police")))
```

```
##           probarr probconv  probsen  avgsen log.police
## density 0.07260985 -0.227912 -0.3005332 0.0715956 0.3282668
```

The correlation matrix confirms our assumptions.

Since *density* is positively skewed, it will benefit with a log transformation because its values are non-zero and positive.

```
# before and after log transformation
hist(data$density); hist(log(data$density))
```



Although *density* has a direct relationship with *urban*, we decided not to include *urban* because there are other factors, such as wage discrepancy between the wealthy and the poor that can potentially exist in urban counties which our current dataset does not have. In other words, by including *urban*, we might fall prey to

the omitted variable bias because there are potentially multiple variables that influence *urban* which are not available in our current dataset.

We chose tax revenue per capita, *tax*, on a similar basis as *density* in that higher tax revenue usually equates to more funding for protection services and therefore lowers the rate of crime. *tax* also has a negative relationship with *crime* and a negative relationship with most of the explanatory variables in Model 1 except for *police*, which we assume will reflect in its coefficients being positive. Again, our reasonings are similar to *density* in that the more money a county has to pay for protection services, such as police, the less likely an individual will commit crime and therefore the lower the probability of arrest, conviction and sentencing is in that county.


```
# list the correlation between tax and model 1's explanatory variables
cor(subset(data, select="tax"),
    subset(data, select=c("probarr", "probconv", "probsen", "avgsen", "log.police")))
```

```
##          probarr  probconv  probsen  avgsen log.police
## tax -0.09236051 -0.1273896 -0.137191 0.08654323 0.4009476
```

The correlatin matrix again confirms our assumptions.

*tax* can also benefit with a log transformation because its distribution is skewed and its values are non-zero and positive.

```
# before and after log transformation
hist(data$tax); hist(log(data$tax))
```



Lastly, we chose *mix* as an indirect effect to *crime* because we are under the impression that the *mix* variable which reflects the ratio of face-to-face crime over all other crimes is a good indicator of violent crimes, and violent crimes have a direct impact to the average sentence a criminal will receive. We understand that it most likely has an effect on *probsen* and *avgsen* but we do not know what type of relationship it has. For this reason, we list the correlation between *mix* and the explanatory variables of Model 1 to get a better understanding.

```
# list the correlation between mix and model 1's explanatory variables
cor(subset(data, select="mix"),
    subset(data, select=c("probarr", "probconv", "probsen", "avgsen", "log.police")))
```

```
##          probarr  probconv  probsen  avgsen log.police
## mix 0.1165888 -0.3042512 0.412898 -0.141705 0.03922583
```

Although *mix* is positively skewed as well, we decided against taking the log transformation of *mix* because it can makes its values between 0 and 1 more extreme.

Before we propose our second model, we would like to discuss the reasons why we decided not to include the other variables as covariates.

Neighborhood plays a central role in fostering the tendency of a person committing a crime. Although we are given geographic locations such as *west* and *central* and neither *west* nor *central*, it does not inform us whether those counties are considered safe or unsafe. Knowing the unsafe rating of a county can give us a better picture of crime rates in those neighborhoods. Also, just having geographic locations do not inform us about laws enacted for safety in those particular regions. We will never know, for example, if counties in western region enact stricter laws than those in the central region. For this reason, we did not consider *west* and *central* as variables in Model 2.

In addition, including *ymale* and *pctmin* will introduce omitted variable bias in Model 2 because typically a county that has low education, high percentage of young male, and high percentage of minority will induce a high crime rate. Education plays a critical role when taking into consideration *ymale* and *pctmin*, and including *ymale* and *pctmin* without the education variable will open up to bias in the model.

Although wage is a good determinant of crime because those who are poor have a higher propensity to commit crime out of financial needs, the groupings of the wage variables do not provide us insight as to the different financial groups between the poor, the middle class, and the wealthy. For example, service industry is mostly thought of as jobs with high number of minimum wage workers, but as the outlier points out in *wageser*, it might be possible for someone to work in the service industry and earn a well-off paycheck if they worked, for example, in a five star hotel. Transportation industry can also be considered as mostly jobs with high number of minimum wage workers, but again, a pilot works in the transportation industry and gets paid well. It is for this reason that we did not include the wage variables in the model.

Hence, we can propose our second model as follows with all the indirect variables included:

$$\log(\text{crime}) = \beta_0 + \beta_1 \cdot \text{probarr} + \beta_2 \cdot \text{probconv} + \beta_3 \cdot \text{probsen} + \beta_4 \cdot \text{avgsen} + \beta_5 \cdot \log(\text{police}) + \beta_6 \cdot \log(\text{density}) + \beta_7 \cdot \log(\text{tax}) + \beta_8 \cdot \text{mix} + u$$

```
m2 <- lm(log(crime) ~ probarr + probconv + probsen + avgsen + log(police) +
         log(density) + log(tax) + mix, data=data)
```

## CLM 1 - 6

1. The model is specified such that the dependent variable is a linear function of the explanatory variables.
2. As discussed in Model 1, we do not know how the survey is collected, but we assume that the variables are representative of the entire population distribution.
3. The variables are not perfectly correlated and the VIFs are low, so there is no perfect multicollinearity of the independent variables.

```
# correlation matrix of explanatory variables
data$log.density <- log(data$density)
data$log.tax <- log(data$tax)
cor(data.matrix(
  subset(data, select=c("probarr", "probconv", "probsen", "avgsen", "log.police",
    "log.density", "log.tax", "mix"))))
```

```
##           probarr      probconv      probsen      avgsen      log.police
## probarr      1.00000000  0.011022645  0.04583324 -0.09468083  0.010413494
## probconv      0.01102265  1.000000000 -0.05579621  0.15585232 -0.007574593
## probsen      0.04583324 -0.055796206  1.00000000  0.17869425  0.216243619
## avgsen      -0.09468083  0.155852319  0.17869425  1.00000000  0.437293263
## log.police    0.01041349 -0.007574593  0.21624362  0.43729326  1.000000000
## log.density   0.07903876 -0.234628931 -0.38514738  0.07458756  0.289033366
## log.tax      -0.06462896 -0.132677900 -0.13078760  0.06101460  0.382987239
## mix           0.11658883 -0.304251245  0.41289804 -0.14170497  0.039225827
##           log.density      log.tax      mix
```

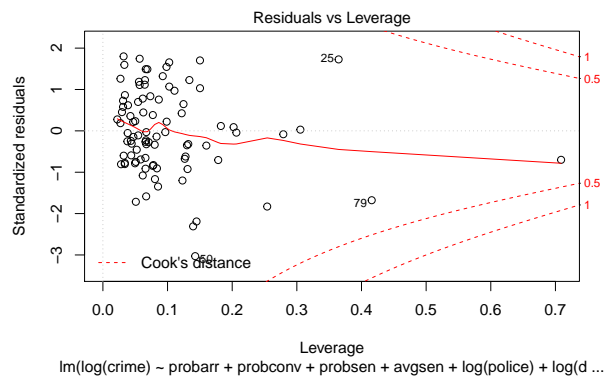

```
# verify VIFs are less than 10
vif(m2)
```

4. Zero-conditional mean assumption holds because the spline line remains close to zero in the residual vs fitted plot, there is no outliers that have high influence, and the covariances are very close to zero indicating the likelihood of being exogenous.

```
lapply(subset(data, select=c("probarr", "probconv", "probsen", "avgsen", "log.police", "log.density", "log.density2")),
       function(var) cov(var, m1$residuals))
```

16





5. The assumption of homoskedasticity is met because the residuals appear randomly spread in the scale-location plot, and the p-value is statistically not significant in both Breusch-pagan and Score-test.

```
# plot scale-location plot
```

```
plot(m1, which=3)
```

```
# Breusch-pagan test
```

```
bptest(m1)
```

```
##
```

```
## studentized Breusch-Pagan test
```

```
##
```

```
## data: m1
```

```
## BP = 6.1759, df = 5, p-value = 0.2895
```


```
# Score-test for non-constant error variance
```

```
ncvTest(m1)
```

```
## Non-constant Variance Score Test
```

```
## Variance formula: ~ fitted.values
```

```
## Chisquare = 1.496155 Df = 1 p = 0.2212639
```



6. Although the assumption is not met, given the substantial sample size, we can be confident that due to OLS asymptotics the distribution of the residuals will be approximately normal.

```
# normality of standard residuals
```

```
rstnd = rstandard(m2)
```

```
hist(rstnd, main="Histogram standard residuals", breaks=50, freq=FALSE)
```

```
curve(dnorm(x, mean=0, sd=sd(rstnd)), col="red", lwd=2, add=TRUE)
```

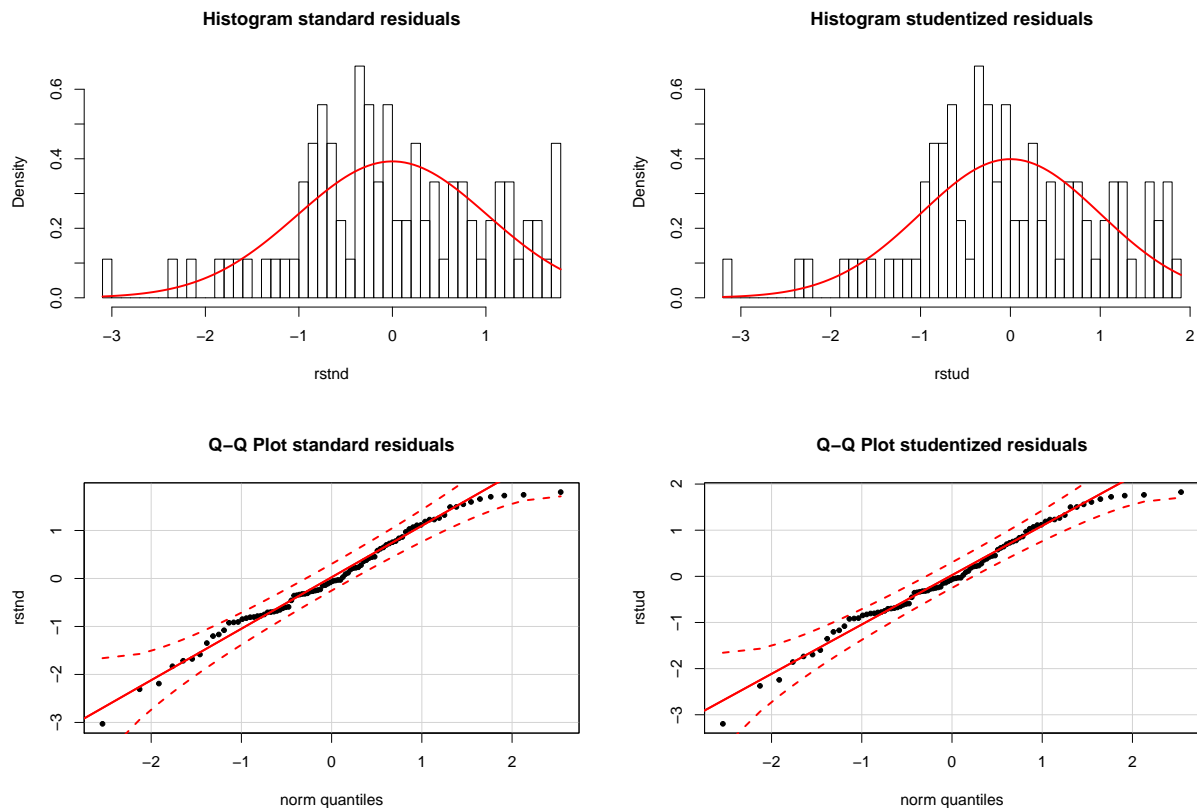
```
# normality of studentized residuals
```

```

rstud = rstudent(m2)
hist(rstud, main="Histogram studentized residuals", breaks=50, freq=FALSE)
curve(dnorm(x, mean=0, sd=1), col="red", lwd=2, add=TRUE)
# Q-Q plot standard residuals
qqPlot(rstnd, distribution="norm", pch=20, main="Q-Q Plot standard residuals")
qqline(rstnd, col="red", lwd=2)
# Q-Q plot studentized residuals
qqPlot(rstud, distribution="norm", pch=20, main="Q-Q Plot studentized residuals")
qqline(rstud, col="red", lwd=2)
# check sample size for model 2
nobs(m2)

```

```
## [1] 90
```



Since all six assumptions of the Classical Linear Model are met, we can assume that the OLS estimators are consistent, normally distributed and BLUE.


## Model 3: most, if not all, other covariates

Adding arbitrary variables uses up valuable degrees of freedom which increases the standard error which in turn reduces the precision of this estimate. This effect is especially pronounced in small sample size.

```

# Look into wageser with an unusual max of 2177.1
boxplot(data$wageser)

```



2177.1 is clearly an outlier in the data and could be possibly due to a measurement error. In this case, we decide not to include this *wageser* to avoid confounding bias caused by potentially inaccurate data.

Moreover, let's look at how *west* and *central* distribute in the sample

```
unique(cbind(data$west, data$central))
```

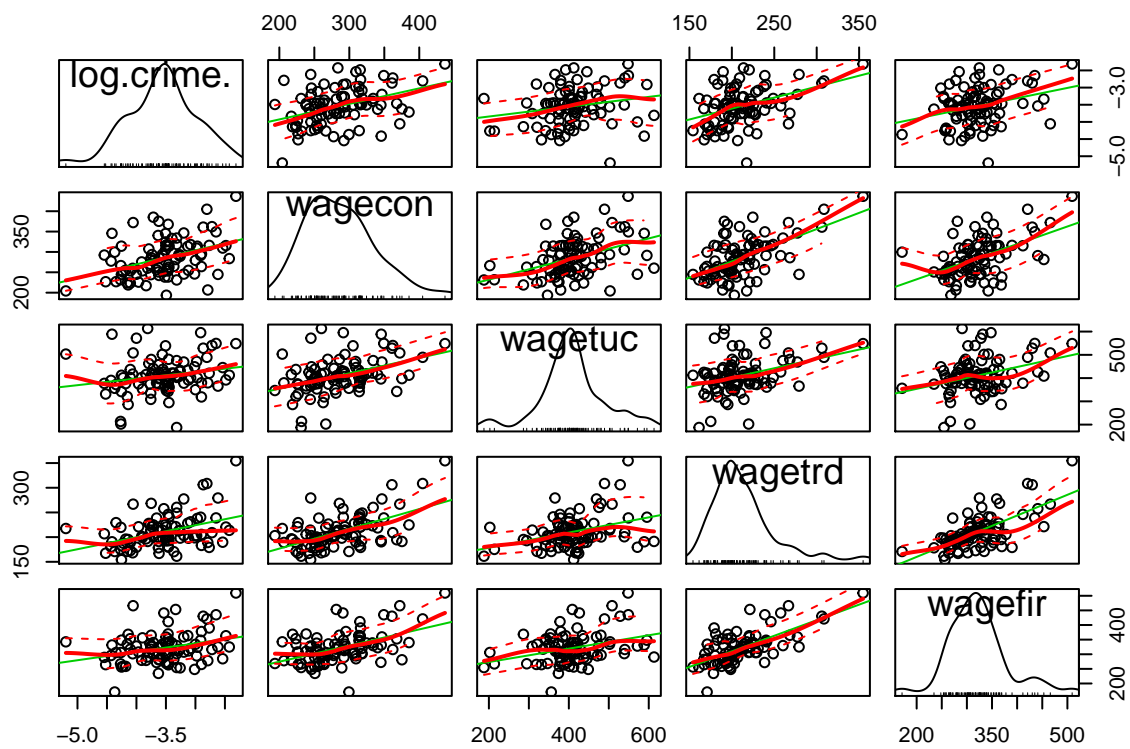
```
##      [,1] [,2]
## [1,]    1    0
## [2,]    0    1
## [3,]    0    0
```

Note that although we don't have any counties that are both in west and central, as expected, we see some counties that neither in west or central. In order to consider the effect on different regions, we will need to use both indicator variables in our model.

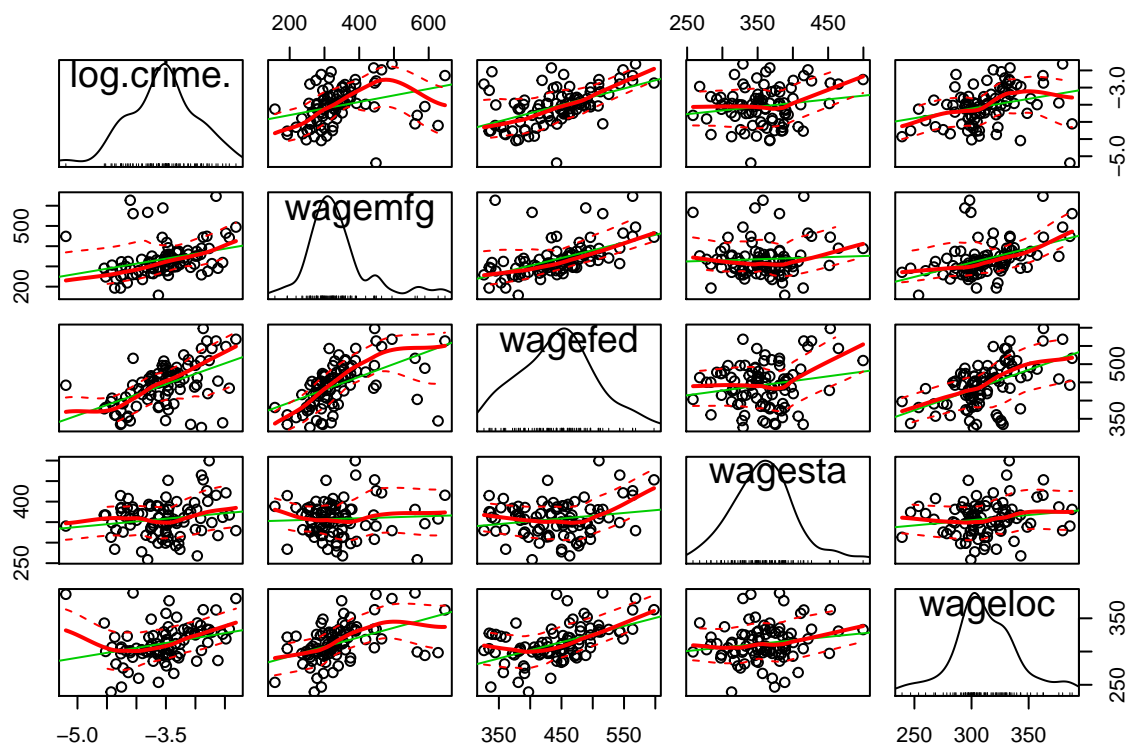
Out of all the variables that aren't included in Model 1 or Model 2, we have decided to include: 1) *west* & *central*, as indicator variables, to control for the regional effect on crime rate 2) *urban*, as an indicator variable, to control for the non-density impact of urbanization on crime rate 3) *wagecon*, *wagetuc*, *wagetrd*, *wagefir*, *wageser*, *wagemfg*, *wagefed*, *wagesta*, *wageloc* to control for the effects of wages in different industries have on our crime rate 4) *pctmin* and *ymale* to control for the demographic effect on crime rate

Now let's look the relationship between the selected variables and our dependent variable *log(crime)*


```
scatterplotMatrix(~ log(crime) + wagecon + wagetuc + wagetrd + wagefir, data = data)
```



```
scatterplotMatrix(~ log(crime) + wagemfg + wagefed + wagesta + wageloc, data = data)
```



```
scatterplotMatrix(~ log(crime) + pctmin + ymale, data = data)
```




None of the variables shows strong evidence of non-linear relationship with the dependent variable  $\log(\text{crime})$ .

We notice that the distributions for the variables  $\text{wagecon}$ ,  $\text{wagetrd}$ ,  $\text{wagefir}$ ,  $\text{wagemfg}$ ,  $\text{wagesta}$ , and  $\text{wageloc}$  are positively skewed, since they are all positive values, we can apply log transformations to all of them. For the ease of the interpretation of the model, we apply log transformation to the other wage related variable  $\text{wagetuc}$  as well.

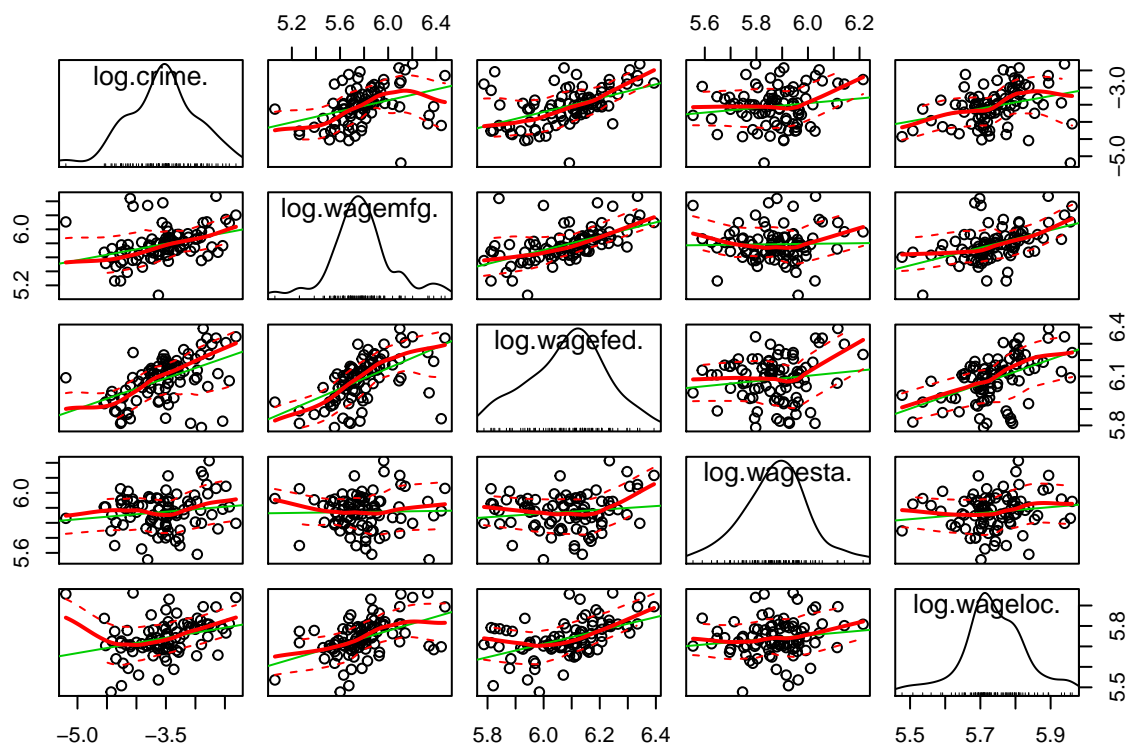
Although  $\text{ymale}$  is also positively skewed and could benefit from a log transformation in terms of normality, it's hard to interpret the slope parameter of its log transformation. Hence, we decided to leave it as it is.

Let's double check the linearity of the relationship between our transformed variables and our dependent variable  $\log(\text{crime})$

```
scatterplotMatrix(~ log(crime) + log(wagecon) + log(wagetuc) + log(wagetrd) + log(wagefir), data = data)
```

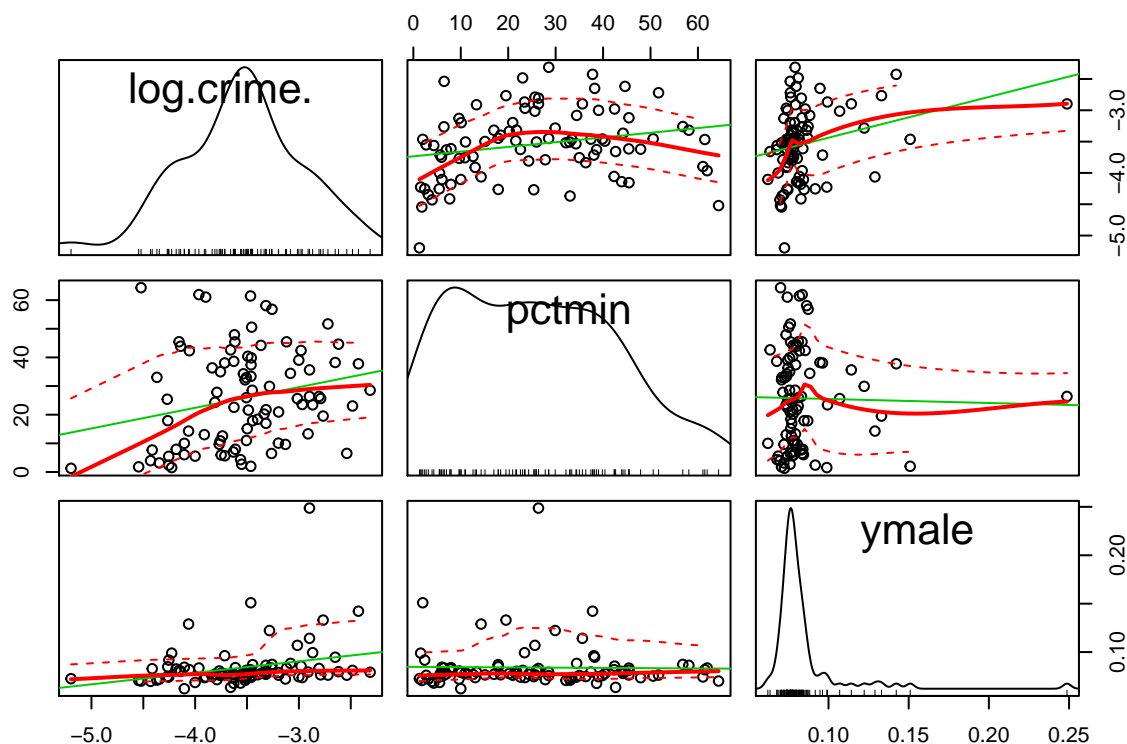


```
scatterplotMatrix(~ log(crime) + log(wagemfg) + log(wagefed) + log(wagesta) + log(wageloc), data = data)
```



```
scatterplotMatrix(~ log(crime) + pctmin + ymale, data = data)
```





We didn't see any strong violation against the linearity of the relationships. Hence, we will propose the following model:

```
m3 <- lm(log(crime) ~ probarr + probconv + probsen + avgsgen + log(police) + log(density) + log(tax) + m
```

## CLM 1 - 6

1. The model is specified such that the dependent variable is a linear function of the explanatory variables.
2. As discussed in Model 1, we do not know how the survey is collected, but we assume that the variables are representative of the entire population distribution.
3. The variables are not perfectly correlated and the VIFs are low, so there is no perfect multicollinearity of the independent variables.

```
# correlation matrix of explanatory variables
data$log.density <- log(data$density)
data$log.tax <- log(data$tax)
cor(data.matrix(
  subset(data, select=c("probarr", "probconv", "probsen", "avgsgen", "log.police",
    "log.density", "log.tax", "mix"))))
```


```
##          probarr    probconv    probsen    avgsgen    log.police
## probarr      1.0000000    0.011022645    0.04583324   -0.09468083    0.010413494
## probconv      0.01102265    1.000000000   -0.05579621    0.15585232   -0.007574593
## probsen       0.04583324   -0.055796206    1.00000000    0.17869425    0.216243619
## avgsgen      -0.09468083    0.155852319    0.17869425    1.00000000    0.437293263
## log.police     0.01041349   -0.007574593    0.21624362    0.43729326    1.000000000
```

```
# verify VIFs are less than 10
vif(m2)
```

4. Zero-conditional mean assumption holds because the spline line remains close to zero in the residual vs fitted plot, there is no outliers that have high influence, and the covariances are very close to zero indicating the likelihood of being exogenous.

```
lapply(subset(data, select=c("probarr", "probconv", "probsen", "avgsen", "log.police", "log.density", "log.employment")),
       function(var) cov(var, m1$residuals))
```

26



5. The assumption of homoskedasticity is met because the residuals appear randomly spread in the scale-location plot, and the p-value is statistically not significant in both Breusch-pagan and Score-test.

```
# plot scale-location plot
```

```
plot(m1, which=3)
```

```
# Breusch-pagan test
```

```
bptest(m1)
```

```
##
```

```
## studentized Breusch-Pagan test
```

```
##
```

```
## data: m1
```

```
## BP = 6.1759, df = 5, p-value = 0.2895
```


```
# Score-test for non-constant error variance
```

```
ncvTest(m1)
```

```
## Non-constant Variance Score Test
```

```
## Variance formula: ~ fitted.values
```

```
## Chisquare = 1.496155 Df = 1 p = 0.2212639
```



6. Although the assumption is not met, given the substantial sample size, we can be confident that due to OLS asymptotics the distribution of the residuals will be approximately normal.

```
# normality of standard residuals
```

```
rstnd = rstandard(m2)
```

```
hist(rstnd, main="Histogram standard residuals", breaks=50, freq=FALSE)
```

```
curve(dnorm(x, mean=0, sd=sd(rstnd)), col="red", lwd=2, add=TRUE)
```

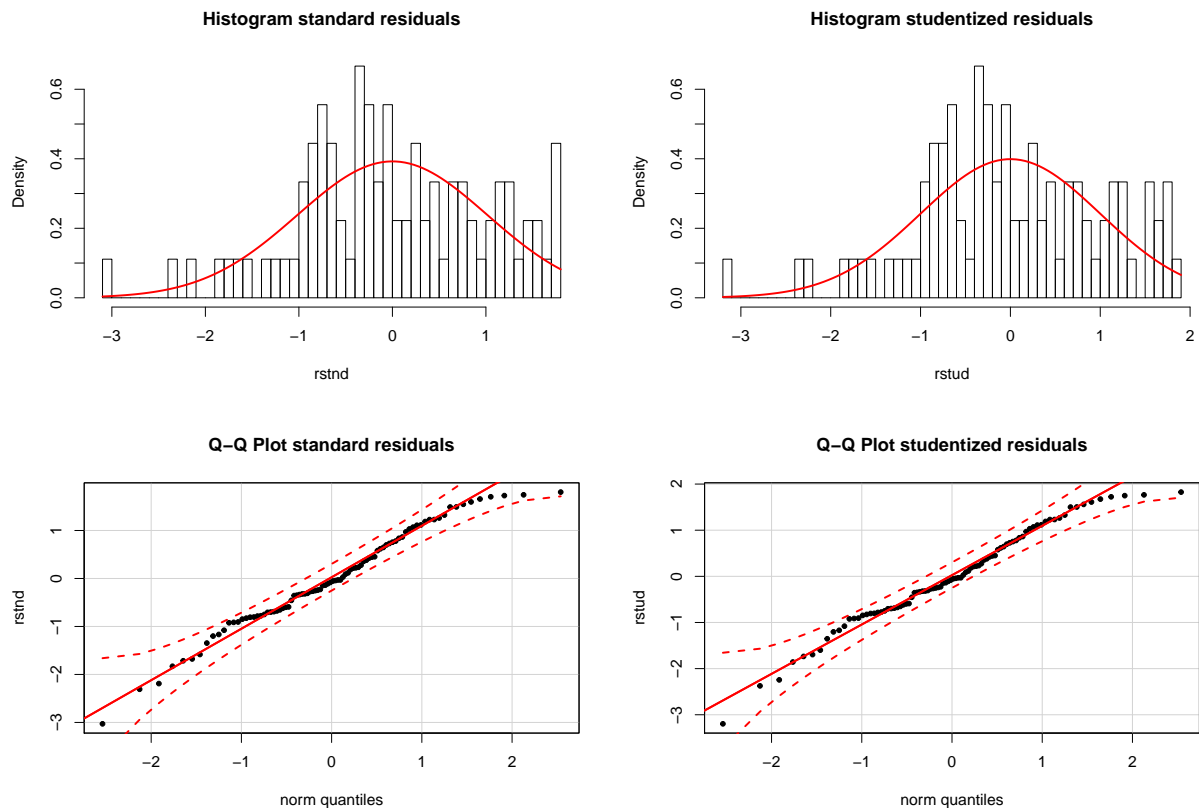
```
# normality of studentized residuals
```

```


rstud = rstudent(m2)
hist(rstud, main="Histogram studentized residuals", breaks=50, freq=FALSE)
curve(dnorm(x, mean=0, sd=1), col="red", lwd=2, add=TRUE)
# Q-Q plot standard residuals
qqPlot(rstnd, distribution="norm", pch=20, main="Q-Q Plot standard residuals")
qqline(rstnd, col="red", lwd=2)
# Q-Q plot studentized residuals
qqPlot(rstud, distribution="norm", pch=20, main="Q-Q Plot studentized residuals")
qqline(rstud, col="red", lwd=2)
# check sample size for model 2
nobs(m2)

```


```
## [1] 90
```



```
plot(m3)
```





Fitted values  
 $\ln(\log(\text{crime})) \sim \text{probarr} + \text{probconv} + \text{probsen} + \text{avg sen} + \log(\text{police}) + \log(\text{d} \dots$



Theoretical Quantiles

$\ln(\log(\text{crime})) \sim \text{probarr} + \text{probconv} + \text{probsen} + \text{avgsen} + \log(\text{police}) + \log(d \dots$





$\text{lm}(\log(\text{crime}) \sim \text{probarr} + \text{probconv} + \text{probsen} + \text{avgsen} + \log(\text{police}) + \log(\text{density}))$

```
bptest(m3)
```

```
##
## studentized Breusch-Pagan test
##
## data: m3
## BP = 41.375, df = 20, p-value = 0.003333
```

Similar to Model 1, Model 3 shows 1) no strong violation of zero-conditional mean, according to the Residuals vs Fitted plot, 2) has no outlier with large influence and 3) shows skewedness in the QQ plot for the residuals. However, Model 3 shows a violation of homoskedasticity according to both the Residuals vs Fitted and the Scale-Location plots.

In response, just like what we normally would have done to be conservative, we will use the heteroskedasticity-robust standard errors for the hypothesis tests on the slope parameters.

One typical concern when including variables that tend to be highly correlated, like wages, is that the multi-collinearity will inflate the variance of the OLS estimated parameters. And to investigate this, we compute the variance inflation factors (vif) of each variable in the model:

```
vif(m3)
```

```
##      probarr      probconv      probsen      avgsen      log(police)
##      1.194093      1.622417      1.893087      1.659673      2.581225
## log(density)      log(tax)          mix          west          central
##      4.776603      2.074658      2.042717      1.733676      1.572025
##      urban log(wagecon) log(wagetuc) log(wagetrd) log(wagefir)
##      2.569216      2.135744      1.707301      3.060257      2.596568
## log(wagemfg) log(wagefed) log(wagesta) log(wageloc)          ymale
```



```
##      2.262547      3.045837      1.653602      2.258130      1.536937
```

Looking at the variance inflation factors, we don't see any variable with high ( $>10$ ) vif, meaning that we don't have high multi-collinearity in our model.

## Summary of Models

```
stargazer(m1, m2, m3, type="latex")
title="Linear Models Predicting Log of Crime Rate"
se = list(se.m1, se.m2, se.m3, omit.stat=c("f", "ser"),
          star.cutoffs=c(0.05, 0.01, 0.001))
```

Next, we can use the heteroskedasticity-robust standard errors to perform t-tests on our estimated parameters:

```
#compute standard errors for model 3 for regression table
(se.m3 = sqrt(diag(vcovHC(m3))))
```

```
## (Intercept)      probarr      probconv      probsen      avgsen
##  4.56536816  0.39500400  0.13179399  0.32653977  0.01551738
## log(police) log(density)      log(tax)      mix      west
##  0.19121534  0.06930902  0.21640558  0.72062234  0.07054323
##      central      urban log(wagecon) log(wagetuc) log(wagetrd)
##  0.08158125  0.14418969  0.24055392  0.28746455  0.41884120
## log(wagefir) log(wagemfg) log(wagefed) log(wagesta) log(wageloc)
##  0.31970564  0.16339323  0.45281284  0.32916485  0.66711837
##      ymale
##  1.78928212
```

```
# We pass the standard errors into stargazer through
# the se argument.
stargazer(m3, type = "text", omit.stat = "f",
          se = list(se.m3),
          star.cutoffs = c(0.05, 0.01, 0.001))
```

```
##
## =====
##                               Dependent variable:
##                               -----
##                               log(crime)
## -----
## probarr                      0.101
##                               (0.395)
##
## probconv                     -0.523***
##                               (0.132)
##
## probsen                      -1.456***
##                               (0.327)
##
## avgsen                       -0.024
##                               (0.016)
##
## log(police)                  0.490*
##                               (0.191)
##
```

```

## log(density)          0.184**
##                      (0.069)
##
## log(tax)              -0.070
##                      (0.216)
##
## mix                   0.028
##                      (0.721)
##
## west                  -0.294***
##                      (0.071)
##
## central               -0.515***
##                      (0.082)
##
## urban                 0.139
##                      (0.144)
##
## log(wagecon)          0.144
##                      (0.241)
##
## log(wagetuc)          0.113
##                      (0.287)
##
## log(wagetrd)          0.107
##                      (0.419)
##
## log(wagefir)          -0.192
##                      (0.320)
##
## log(wagemfg)          -0.032
##                      (0.163)
##
## log(wagefed)          0.589
##                      (0.453)
##
## log(wagesta)          -0.162
##                      (0.329)
##
## log(wageloc)          -0.080
##                      (0.667)
##
## ymale                 0.998
##                      (1.789)
##
## Constant              -2.030
##                      (4.565)
##
## -----
## Observations          90
## R2                    0.854
## Adjusted R2           0.811
## Residual Std. Error   0.238 (df = 69)
## =====

```

## Note:                      \*p<0.05; \*\*p<0.01; \*\*\*p<0.001

According to the hypothesis tests using the heteroskedasticity-robust standard errors, we can see that variables *probconv*, *probsen*, *west*, *central* show strong statistical significance (p-value < 0.001), same with *log(density)* (p-value < 0.01) and *log(police)* (p-value < 0.05).

Moreover, to interpret the estimated parameters for: 1) *probconv*: keeping all other variables constant, a 0.1 increase in *probconv* results in a 5.23% decrease in crime rate 2) *probsen*: keeping all other variables constant, a 0.1 increase in *probsen* results in a 14.56% decrease in crime rate 3) *west*: keeping all other variables constant, being in west region decreases crime rate by 29.4% 4) *central*: keeping all other variables constant, being in central region decreases crime rate by 51.5% 5) *log(density)*: keeping all other variables constant, an 1% increase in density increases crime rate by 0.18% 6) *log(police)*: keeping all other variables constant, an 1% increase in police per capita increases crime rate by 0.49%

And from this, *probconv*, *probsen*, *west*, *central* show strong practical significance.

## Discussion of Causality

## Conclusion

TODO: These are notes for policy suggestions

Read more on <http://lawenforcementleaders.org/wp-content/uploads/2017/10/Law-Enforcement-Letter-to-President-Trump-and-Attorney-General-Sessions.pdf>

this effort is undercut by a diffuse focus. Attorney General Sessions' regular statements encouraging law enforcement to focus on drug and nonviolent offenders divert officers away from that vital mission. Law enforcement resources are limited. Focusing on low-level non-violent offenders means less time to stop and bring to justice the most dangerous offenders.

From our experience, we do not believe that always seeking the longest possible sentence will make our country safer. More than 25 percent of the Justice Department's budget is consumed by federal prisons.[iv] Every unnecessary dollar spent on prisons is a dollar not spent on policing. And often, the best way to prevent recidivism is through treatment, not prison. Responsibly reducing incarceration will free funding and time for our officers to focus on targeting and preventing violent crime, making our streets safer.[v]

## References:

"Shattering"Broken Windows": An Analysis of San Francisco's Alternative Crime Policies", CENTER ON JUVENILE AND CRIMINAL JUSTICE, October 1999 <http://www.cjcj.org/uploads/cjcj/documents/shattering.pdf>

Jackman, Tom. "Nation's top cops, prosecutors urge Trump not to roll back successful crime policies." The Washington Post, WP Company, 18 Oct. 2017, [www.washingtonpost.com/news/true-crime/wp/2017/10/18/nations-top-cops-prosecutors-urge-trump-not-to-roll-back-successful-crime-policies/?utm\\_term=.53fb295eac1e](http://www.washingtonpost.com/news/true-crime/wp/2017/10/18/nations-top-cops-prosecutors-urge-trump-not-to-roll-back-successful-crime-policies/?utm_term=.53fb295eac1e).